STARTER KIT

AI for program managers running a GTM/CX portfolio

How to hand the repetitive work to agents, keep people on the judgment calls, and get non-technical teams to adopt it, while staying security-first.

SECTION	PAGE
1. The portfolio cycle and where agents fit	2
2. The adoption strategy for non-technical teams	3–4
3. Five use cases	5–6
4. Guardrails for a security-first company	7
5. Using the skills and the 90-day path	7

October 2026 · Free to reuse and adapt (CC BY 4.0) · github.com/sundar-1566/pgm-ai-starter-kit · All examples use made-up data.

1 · THE PREMISE

A PgM on a GTM/CX portfolio runs this cycle every quarter

Based on a typical GTM/CX portfolio: portfolio planning, change adoption, and executive alignment, at 50+ projects a quarter.

#	STAGE	WHERE THE TIME GOES	WHAT AN AGENT DOES	FIT
1	Intake	Half-filled requests, duplicates, chasing details	Structures the request, lists missing info, checks overlap	Good fit
2	Evaluate & prioritize	Building and defending each case	Drafts the score with reasons; council decides	Assist
3	Sequence & resource	Spotting collisions across 50+ projects	Flags date and team clashes	Assist
4	Weekly governance	Rebuilding status from five places, every week	Drafts the roll-up from Slack, email, notes, Jira	Good fit
5	Deliver the change	Writing UAT scripts, chasing testers	Drafts UAT scripts when QA passes	Good fit
6	Adoption & friction	Reading scattered field feedback	Groups feedback into friction themes and fixes	Good fit
7	Exec forums	Rebuilding the QBR story each quarter	Builds the pre-read from 13 weekly roll-ups	Good fit
8	Close & learn	Lessons that never cross projects	Finds patterns across retros	Good fit

How to pick the first wins

Start where the work is frequent and mostly text, the inputs already exist (Slack, Jira, notes), the output is only a draft, and a person can check it quickly. Leave anything that writes into quoting, billing, or licensing systems for last.

The math on status alone (example; adjust to your team)

50 projects × 20 minutes a week to collect updates and write status = **about 17 PgM hours a week**, or close to half a person, before intake chasing or QBR prep. If an agent drafts 60% of it, that's **about 10 hours a week** back for risk calls, stakeholder time, and better exec storytelling.

The PMO already owns change management. Make it the engine for AI adoption across GTM/CX.

Engineering teams adopt AI on their own. Sales Ops, Deal Desk, CS Ops, Enablement, and the PMO usually don't. The PMO already sits across all of them, runs the adoption playbook, and reports to leaders on it, which makes it the natural owner.

Where non-technical teams stall

Curious → Trying → Trusting → Relying → Redesigning. Most teams stall between trying it and trusting the output enough to stop redoing it.

WHAT GETS IN THE WAY	WHAT WORKED
The blank prompt box People don't know what to ask, so they ask once and stop.	**Ready-made skills** Recipes for real tasks. They supply the input, and the recipe does the rest.
"Can I trust this?" One bad output and the tool gets dropped.	**Start personal** Each PgM uses it on their own work first. Track the rewrite rate so trust is visible.
"Is this even allowed?" Security uncertainty stops people more often than an actual rule does.	**Make the guardrails visible** Publish what's approved, what agents can touch, and the log.
No time to learn Busy teams won't take a course to save time later.	**Champions and short show-and-tells** Two peers per team run a weekly 15-minute session on real projects.
"Will this replace me?" Quiet fear slows use, especially in ops roles.	**Show where the returned time goes** Report hours returned and what they went to.
Uneven results Ten people write ten prompts and get ten quality levels.	**One shared skill per task** When results are off, fix the recipe instead of coaching each person.

How you know it's working

Hours returned
per PgM per week, before vs. after

Rewrite rate
share of each draft people change; falling means trust is rising

On-time reporting
roll-ups and pre-reads on schedule

Weekly active users
who keeps using it after week 4

Two accelerators that bend the curve

The fixes on the previous page get people to try AI. These two get teams to rely on it.

TRYING → TRUSTING

A shared library of skills and tips, with the best ones rewarded

- One team repository for skills, prompts, and tips. Anyone can add one, and everyone can reuse it.
- Each entry has a named author, the task it covers, and a before-and-after example.
- Recognize the best creations every month, judged by results and how often others reuse them.
- Retire or fix entries that people keep rewriting.

The best ideas come from the people doing the work. A library turns one person's shortcut into the team's standard, and rewarding it makes creating skills part of the job.

TRUSTING → RELYING

Hackathons and lockdown weeks to force the habit

- **Hackathon:** 1–2 days on the team's real pain points. Every pair ships one working skill into the library.
- **Lockdown week:** for one week, a recurring task (weekly status, meeting notes) goes through the agent only, with no manual fallback.
- Leaders take part and use the outputs in their own reviews that week.
- Debrief afterward: what broke, what to fix, what stays.

When the tool stays optional, many people never get past the awkward first weeks. A short, deliberate forcing function gets them through faster than months of encouragement.

Two lessons learned the hard way

LESSON	AVOID	INSTEAD
Token counts don't drive the right behavior	Ranking people or teams by tokens or prompts used. It rewards activity over results, and people learn to burn tokens to look adopted.	Measure outcomes: hours returned, rewrite rate, reports on time, and decisions made faster. Leave token counts off the scorecard.
"The AI did that" is never an excuse	Letting sloppy AI output go out and blaming the tool when it's wrong. Blaming the tool makes leaders doubt every draft that follows.	Whoever sends it owns it. Review every draft before it goes to leaders, the field, or customers, as if you wrote it yourself. The `no-ai-slop` skill helps with that review.

Five places to start

1. Intake triage Portfolio

PAIN	Requests arrive half-filled; some duplicate active work.
AGENT DOES	Structures the request, drafts questions for missing info, checks overlap with active projects, drafts a score with reasons.
EXAMPLE	A one-line Slack ask for a new partner-discount approval step becomes a structured request, two questions for the requester, and a flag that it overlaps an in-flight deal-registration redesign. Recommendation: merge.
MEASURE	Requests complete on first review; duplicates caught before kickoff.

Human sign-off: the portfolio council makes the priority call.

2. Status roll-up → QBR pre-read Portfolio + Exec

PAIN	Status is rebuilt by hand every week from five places.
AGENT DOES	Reads Slack, email, meeting notes, and Jira; matches updates to projects; sets RAG with a reason and a source; flags silent projects; drafts the roll-up. At quarter end, 13 roll-ups become the QBR pre-read.
MEASURE	Minutes per roll-up; rewrite rate under 20% by week 4; on-time delivery.

Human sign-off: PgM approves before anything is posted. The agent can't send email.

3. Meeting → decisions, risks, and draft Jira tickets Portfolio + Exec

PAIN	Every meeting ends with 30 minutes of cleanup, and actions get lost.
AGENT DOES	Separates decided from discussed, lists actions with named owners, updates the risk log, drafts a Jira ticket per action.
MEASURE	Minutes from meeting end to actions logged; tickets created without edits.

Human sign-off: the PgM confirms each ticket (Create, Edit, or Drop) before anything enters Jira.

4. QA pass → UAT scripts (event-triggered) Change delivery

PAIN	Business testers wait days for UAT scripts after QA finishes.
AGENT DOES	Triggered when a story moves to QA Passed. Reads the acceptance criteria and linked docs, writes business-language scripts with a negative test per criterion, checks coverage, creates draft UAT sub-tasks for testers. Stops and asks if the criteria are vague.
MEASURE	Days from QA pass to UAT start (target: same day); criteria covered (target: 100%).

Human sign-off: the business owner approves scripts before testers are notified. The agent can't change ticket status.

5. Adoption friction listener Change

PAIN	You can't tell why people aren't using a new process, or which complaints matter.
AGENT DOES	Reads field Slack, tickets, and survey text after a rollout; groups feedback into six friction types (don't know how, takes too long, doesn't fit my case, don't see why, broken, don't trust it); sizes each; spots workarounds; drafts fixes and a "you said, we did" note.
EXAMPLE	Week 3 of a new quote-approval flow: a hard-to-find field, too many approvals for small discounts, a routing gap for multi-year deals, and a spreadsheet workaround because status isn't visible. Each gets an owner and an effort size.
MEASURE	Time from friction first seen to fix; workarounds retired; adoption trend after each fix cycle.

Human sign-off: the process owner reviews before anything goes to the field. Names are removed; it reports patterns, not people.

Also in the kit: `retro-synthesis`, which finds patterns across project retros and proposes process fixes with owners.

4 · GUARDRAILS

Six ground rules for a security-first company

1. Approved tools only. Use the AI tools and connectors security has approved.

2. Read-only first. Add write access one action at a time.

3. Drafts, not decisions. Agents don't close, approve, send, or change status.

4. A person signs off on anything that reaches leaders, customers, or the field.

5. Log every tool call, including blocked ones, reviewed weekly.

6. Business systems last. Quoting, billing, and licensing wait until trust is earned.

Why a gate and a log come as a pair (a hypothetical). Imagine a routine improvement to an AI agent that quietly gives it access to a new tool. If your launch review only checks which data the agent reads, not which tools it can call, the change sails through. An activity log would catch it: a tool being used where it never was before. The fix is to make tool access its own review gate. A gate prevents the problems you expected, and a log detects the ones you didn't.

5 · GETTING STARTED

Using the skills, and a 90-day path

Each skill is a plain-English recipe (a `SKILL.md` file). No code to install.

- **Claude Code or Desktop:** put each skill folder in `~/.claude/skills/`.
- **Claude Projects or another approved AI workspace:** paste the recipe as project instructions.
- **Any approved chat assistant:** paste the recipe, then your inputs.
- **Try it now:** the `sample-data/` folder has a meeting transcript, a portfolio list, and field feedback.

DAYS	WAVE	WHAT RUNS
1–30	Personal helpers	Meeting → Jira and personal status drafts; two champions; weekly 15-minute show-and-tell
31–60	Shared reporting	Portfolio roll-up from one source; QBR pre-read from the same data; intake triage
61–90	Triggered agents	QA → UAT on a Jira trigger; adoption pulse after every rollout; each new agent clears security review first